

Natural Language Processing with Deep Learning

Neural Networks – a Walkthrough



Navid Rekab-Saz

navid.rekabsaz@jku.at

Agenda

- Introduction
- Non-linearities
- Forward pass & backpropagation
- Softmax & loss function
- Optimization & regularization

Agenda


- **Introduction**
- Non-linearities
- Forward pass & backpropagation
- Softmax & loss function
- Optimization & regularization

Notation

- $a \rightarrow$ scalar
- $\mathbf{b} \rightarrow$ vector
 - i^{th} element of \mathbf{b} is the scalar b_i
- $\mathbf{C} \rightarrow$ matrix
 - i^{th} vector of \mathbf{C} is \mathbf{c}_i
 - j^{th} element of the i^{th} vector of \mathbf{C} is the scalar $c_{i,j}$
- Tensor: generalization of scalar, vector, matrix to any arbitrary dimension

Linear Algebra

Once upon a time in algebra class...



"LIKE WE'LL EVER USE THIS CRAP."

Linear Algebra – Transpose

- a is in $1 \times d$ dimensions $\rightarrow a^T$ is in $d \times 1$ dimensions
- A is in $e \times d$ dimensions $\rightarrow A^T$ is in $d \times e$ dimensions

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}^T = \begin{bmatrix} 1 & 4 \\ 2 & 5 \\ 3 & 6 \end{bmatrix}$$

Linear Algebra – Dot product

- $\mathbf{a} \cdot \mathbf{b}^T = c$

- dimensions: $1 \times d \cdot d \times 1 = 1$

$$\begin{bmatrix} 1 & 2 & 3 \end{bmatrix} \begin{bmatrix} 2 \\ 0 \\ 1 \end{bmatrix} = 5$$

- $\mathbf{a} \cdot \mathbf{B} = \mathbf{c}$

- dimensions: $1 \times d \cdot d \times e = 1 \times e$

$$\begin{bmatrix} 1 & 2 & 3 \end{bmatrix} \begin{bmatrix} 2 & 3 \\ 0 & 1 \\ 1 & -1 \end{bmatrix} = \begin{bmatrix} 5 & 2 \end{bmatrix}$$

- $\mathbf{A} \cdot \mathbf{B} = \mathbf{C}$

- dimensions: $l \times m \cdot m \times n = l \times n$

$$\begin{bmatrix} 1 & 2 & 3 \\ 1 & 0 & 1 \\ 0 & 0 & 5 \\ 4 & 1 & 0 \end{bmatrix} \begin{bmatrix} 2 & 3 \\ 0 & 1 \\ 1 & -1 \end{bmatrix} = \begin{bmatrix} 5 & 2 \\ 3 & 2 \\ 5 & -5 \\ 8 & 13 \end{bmatrix}$$

- **Linear transformation:** dot product of a vector to a matrix

Probability

- Conditional probability

$$p(y|x)$$

- Probability distribution

- For a **discrete** random variable \mathbf{z} with K states
 - $0 \leq p(z_i) \leq 1$
 - $\sum_{i=1}^K p(z_i) = 1$
- E.g. with $K = 4$ states: $[0.2 \quad 0.3 \quad 0.45 \quad 0.05]$

Probability

- Expected value

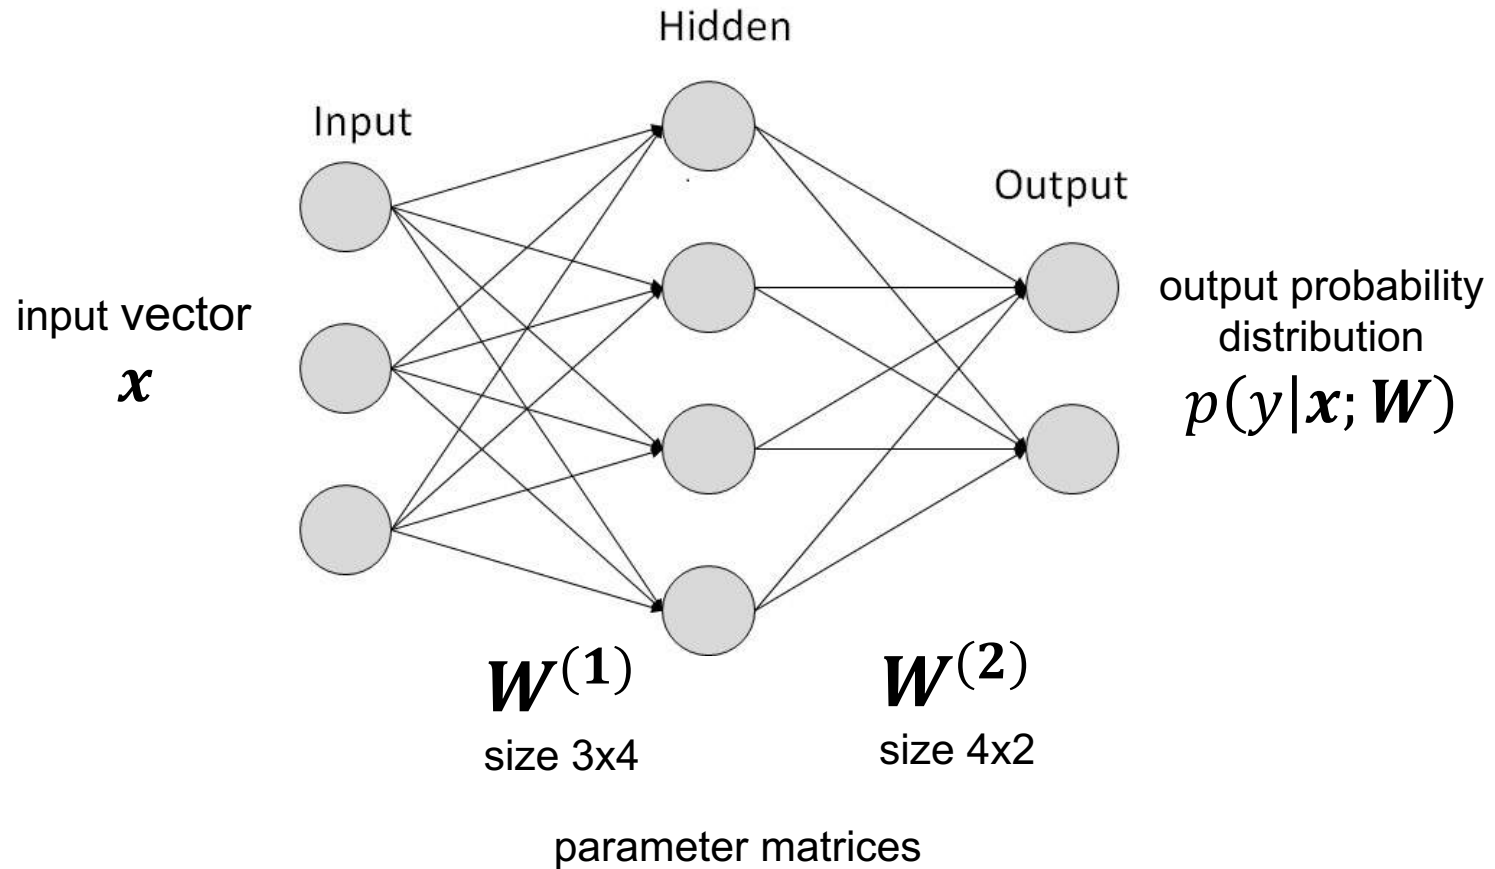
$$\mathbb{E}_{x \sim X}[f] = \frac{1}{|X|} \sum_{x \in X} f(x)$$

- Note: This is an imprecise definition. Though, it suffices for our use in this lecture

Artificial Neural Networks

- Neural Networks are non-linear functions and universal approximators
- They composed of several simple (non-)linear operations
- Neural networks can readily be defined as probabilistic models which estimate $p(y|\mathbf{x}; \mathbf{W})$
 - Given input vector \mathbf{x} and the set of parameters \mathbf{W} , estimate the probability of the output class y

A Feedforward network




Learning with Neural Networks

- Design the network's architecture
- Consider proper **regularization** methods
- Initialize parameters
- Loop until some **exit criteria** are met
 - Sample a **minibatch** from training data \mathcal{D}
 - Loop over data points in the minibatch
 - **Forward pass**: given input \mathbf{x} predict output $p(y|\mathbf{x}; \mathbf{W})$
 - Calculate **loss** function
 - Calculate the **gradient** of each parameter regarding the loss function using the **backpropagation** algorithm
 - **Update** parameters using their gradients


Agenda

- Introduction
- **Non-linearities**
- Forward pass & backpropagation
- Softmax & loss function
- Optimization & regularization

Neural Computation




An Artificial Neuron



Linear


$$f(x) = x$$



Sigmoid

$$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$$


- squashes input between 0 and 1
- Output becomes like a probability value



Hyperbolic Tangent (Tanh)

$$f(x) = \tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$


- squashes input between -1 and 1



Rectified Linear Unit (ReLU)

$$f(x) = \max(0, x)$$

- Good for deep architectures, as it prevents vanishing gradient



Examples

$$\mathbf{x} = [1 \quad 3] \quad \mathbf{W} = \begin{bmatrix} 0.5 & -0.5 & 2 & 0 & 0 \\ 0 & 0 & 0 & 4 & -1 \end{bmatrix}$$

- Linear transformation \mathbf{xW} :

$$\mathbf{xW} = [1 \quad 3] \begin{bmatrix} 0.5 & -0.5 & 2 & 0 & -1 \\ 0 & 0 & 0 & 4 & -1 \end{bmatrix} = [\mathbf{0.5} \quad \mathbf{-0.5} \quad \mathbf{2} \quad \mathbf{12} \quad \mathbf{-4}]$$

- Non-linear transformation $\text{ReLU}(\mathbf{xW})$:

$$\text{ReLU}([0.5 \quad -0.5 \quad 2 \quad 12 \quad -3]) = [\mathbf{0.5} \quad \mathbf{0.0} \quad \mathbf{2} \quad \mathbf{12} \quad \mathbf{0.0}]$$

- Non-linear transformation $\sigma(\mathbf{xW})$:

$$\sigma([0.5 \quad -0.5 \quad 2 \quad 12 \quad -3]) = [\mathbf{0.62} \quad \mathbf{0.37} \quad \mathbf{0.88} \quad \mathbf{0.99} \quad \mathbf{0.018}]$$

- Non-linear transformation $\tanh(\mathbf{xW})$:

$$\tanh([0.5 \quad -0.5 \quad 2 \quad 12 \quad -3]) = [\mathbf{0.46} \quad \mathbf{-0.46} \quad \mathbf{0.96} \quad \mathbf{0.99} \quad \mathbf{-0.99}]$$

Agenda

- Introduction
- Non-linearities
- **Forward pass & backpropagation**
- Softmax & loss function
- Optimization & regularization

Forward pass

- Consider this calculation:

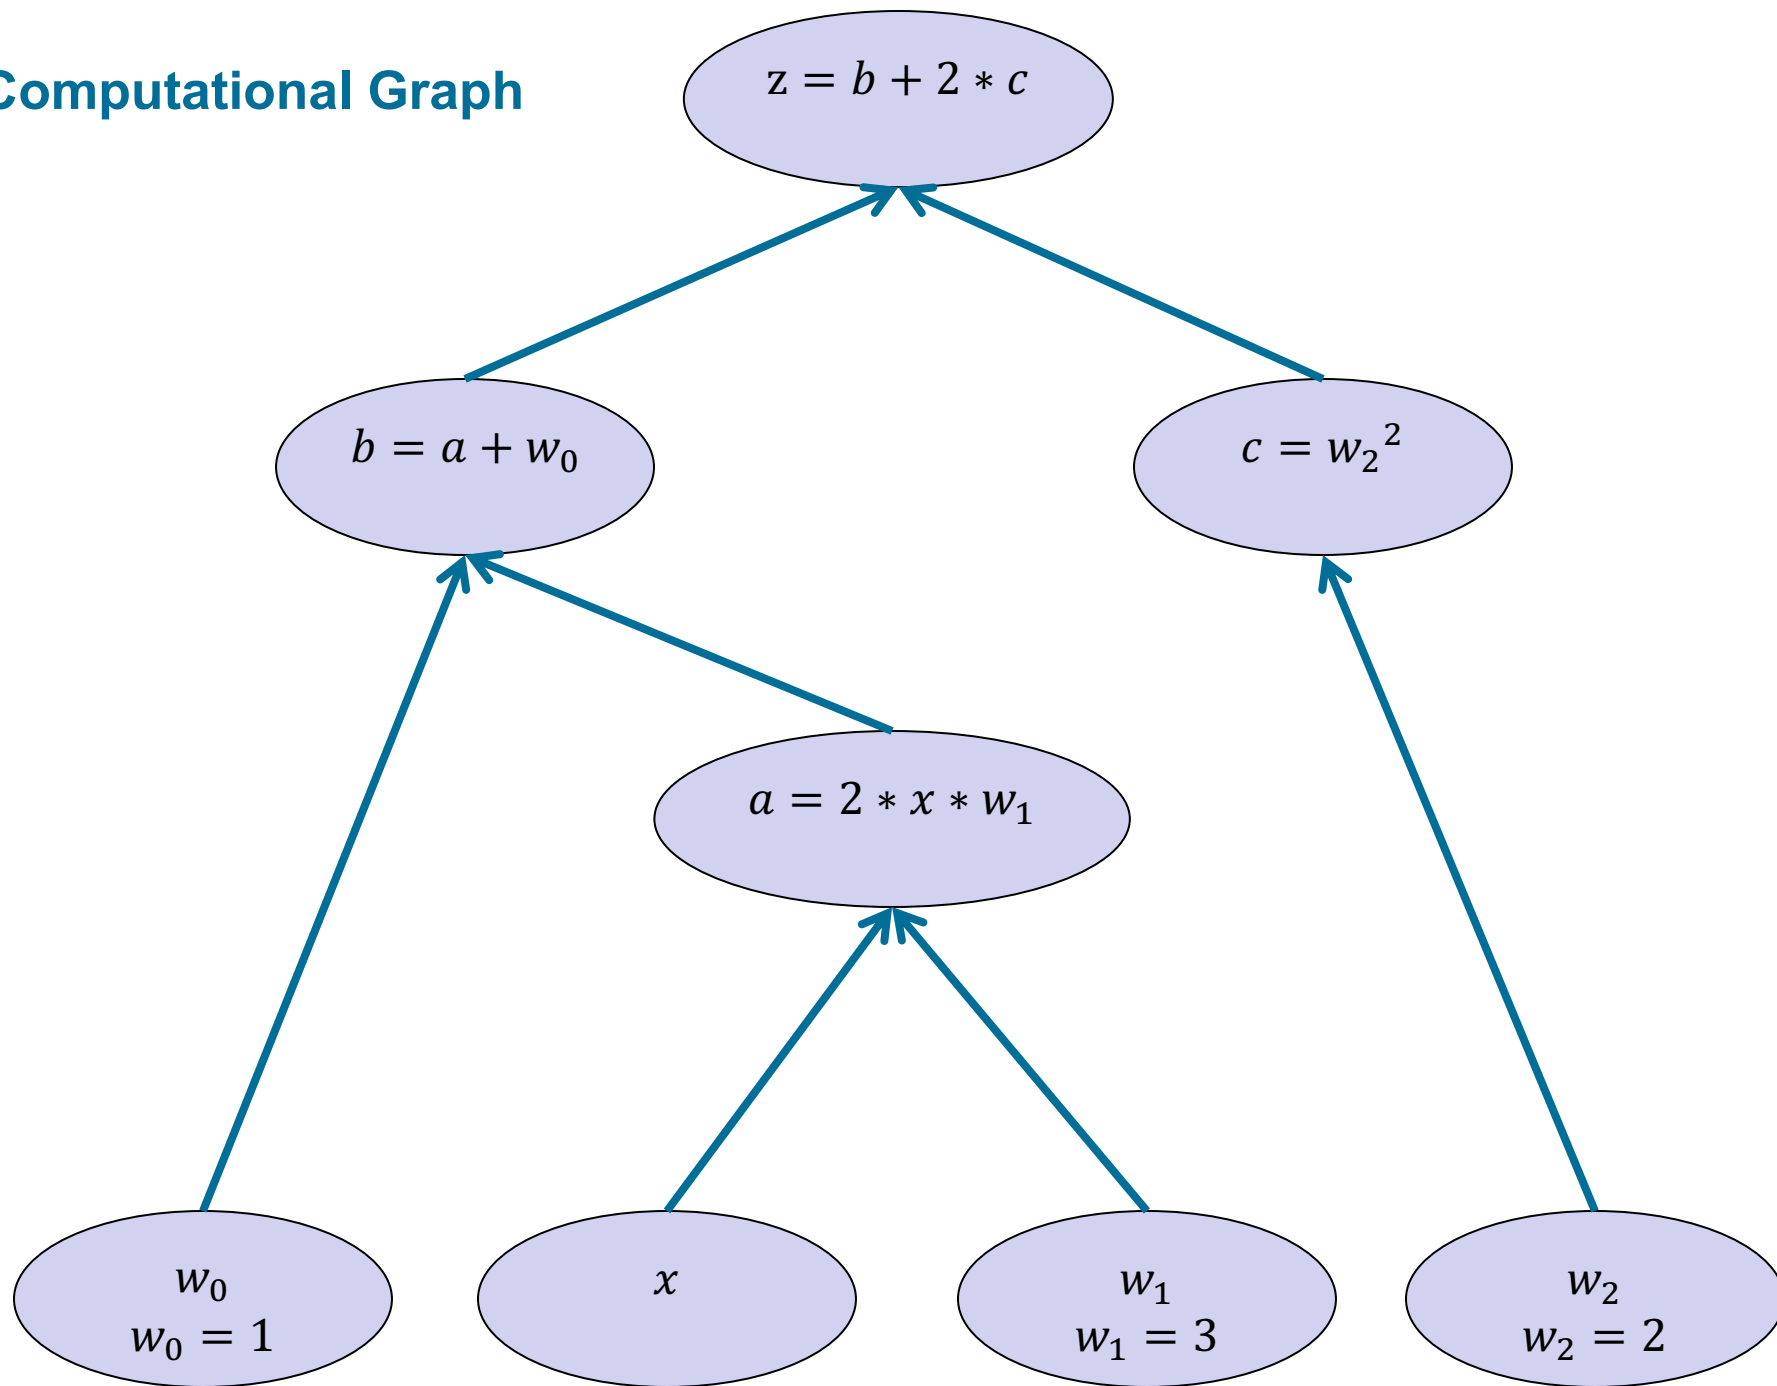
$$z(x; \mathbf{w}) = 2 * w_2^2 + x * w_1 + w_0$$

where x is input and \mathbf{w} is the set of parameters with the initialization
 $w_0 = 1$ $w_1 = 3$ $w_2 = 2$

- Let's break it into intermediary variables:

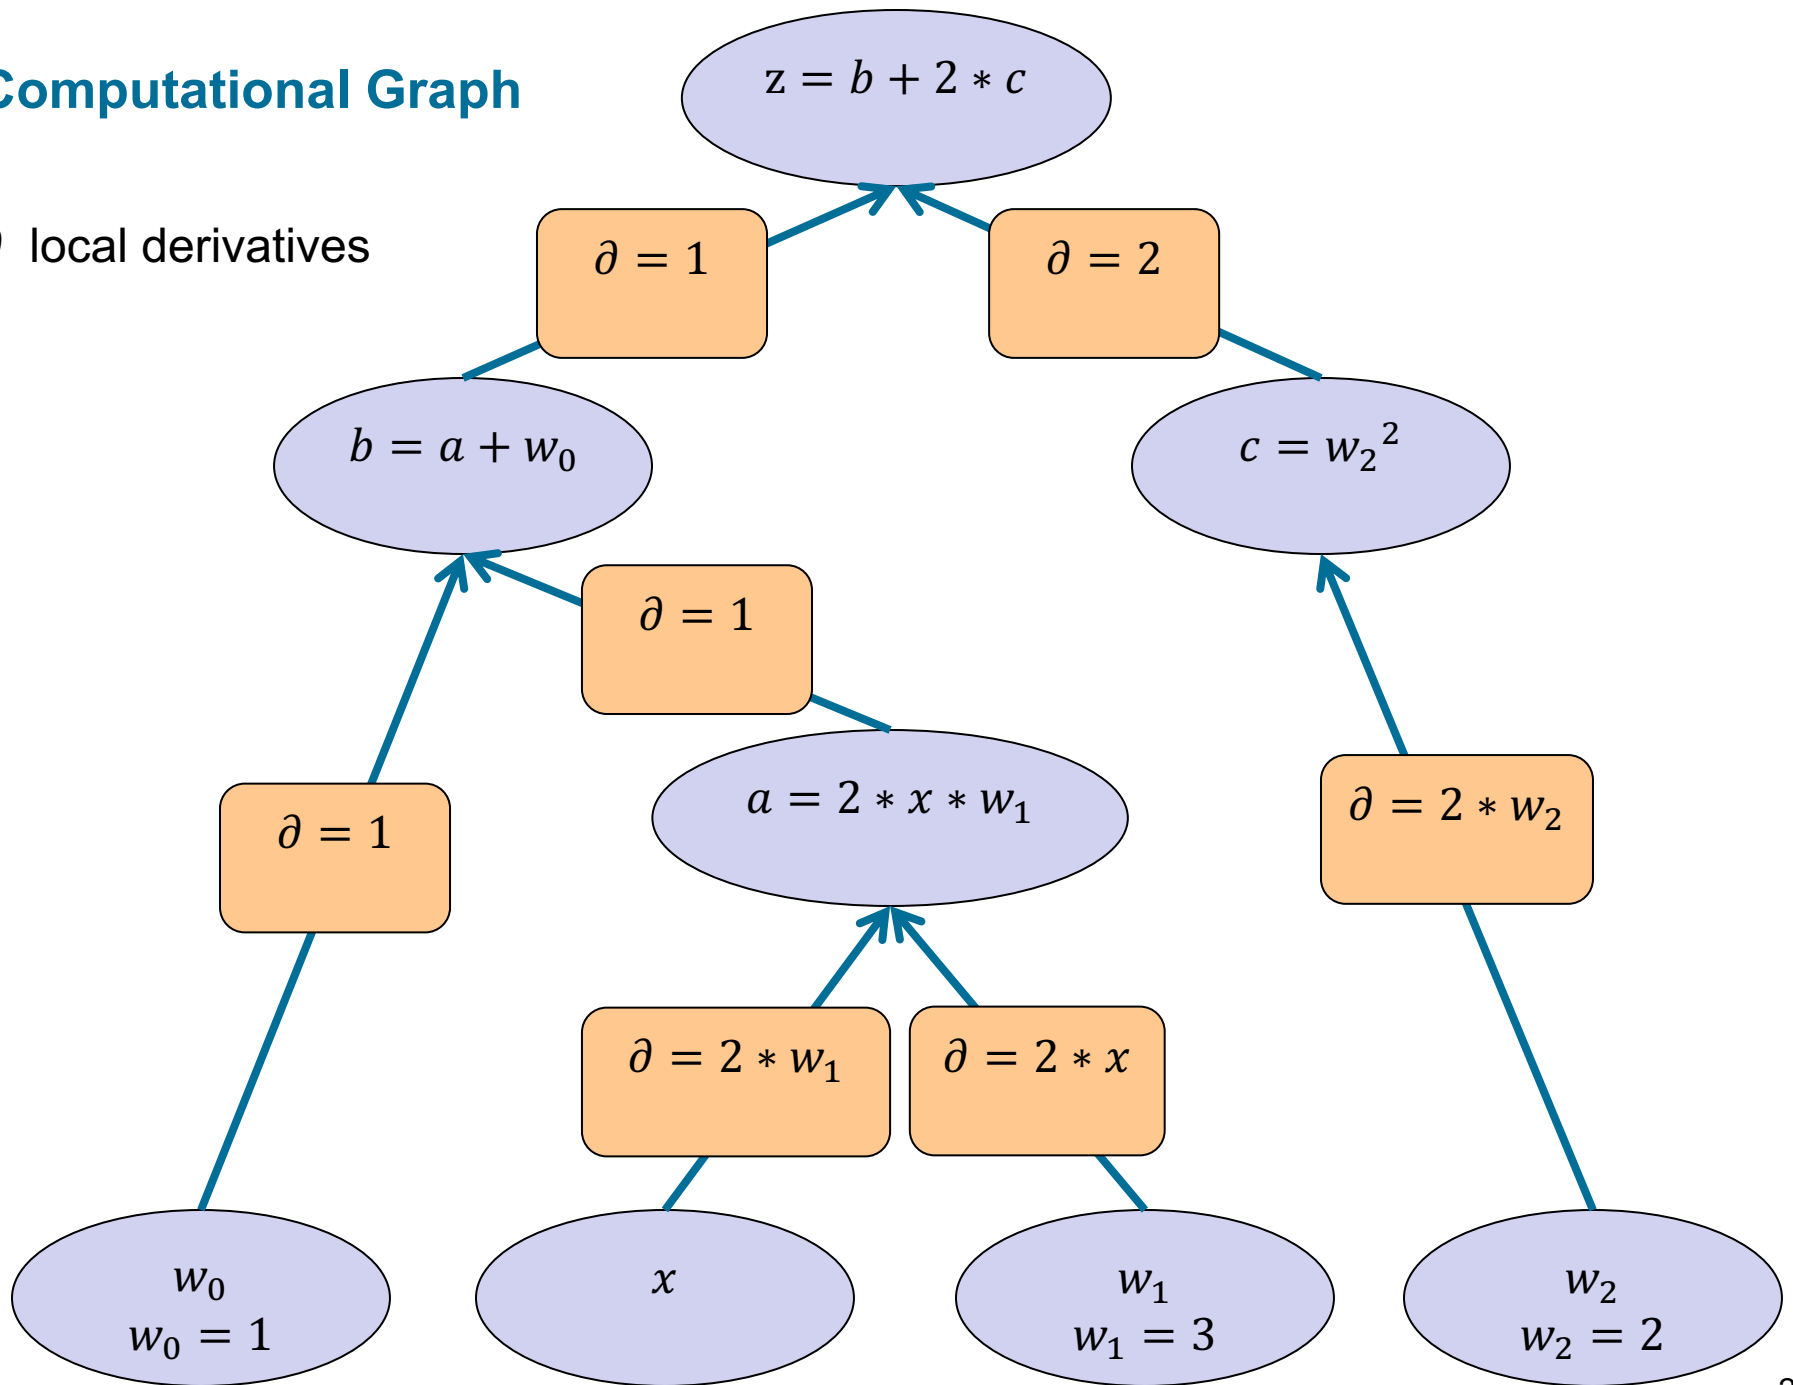
$$\begin{aligned}a &= x * w_1 \\b &= a + w_0 \\c &= w_2^2 \\z &= b + 2 * c\end{aligned}$$

Computational Graph



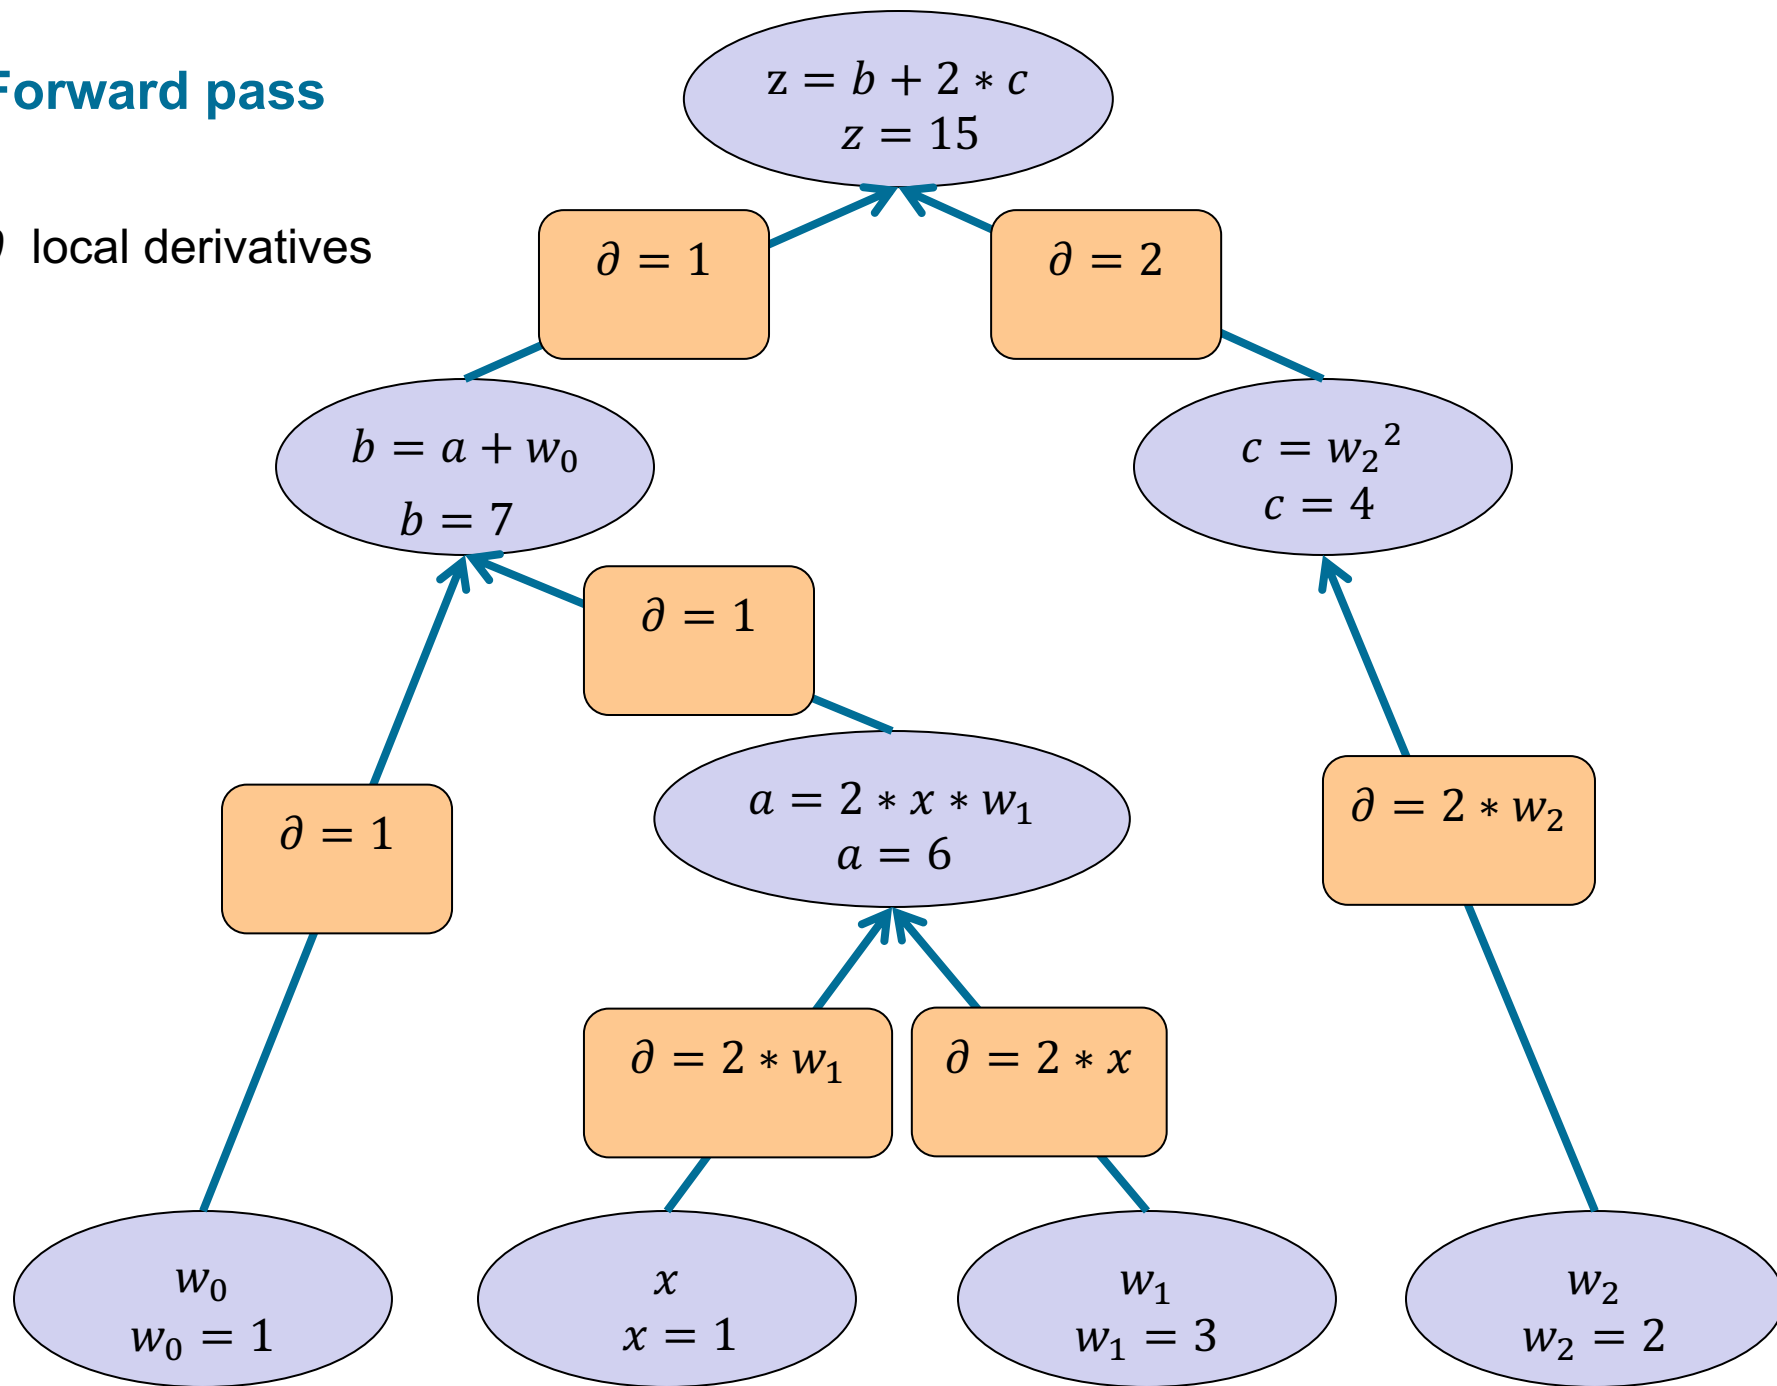
Computational Graph

∂ local derivatives



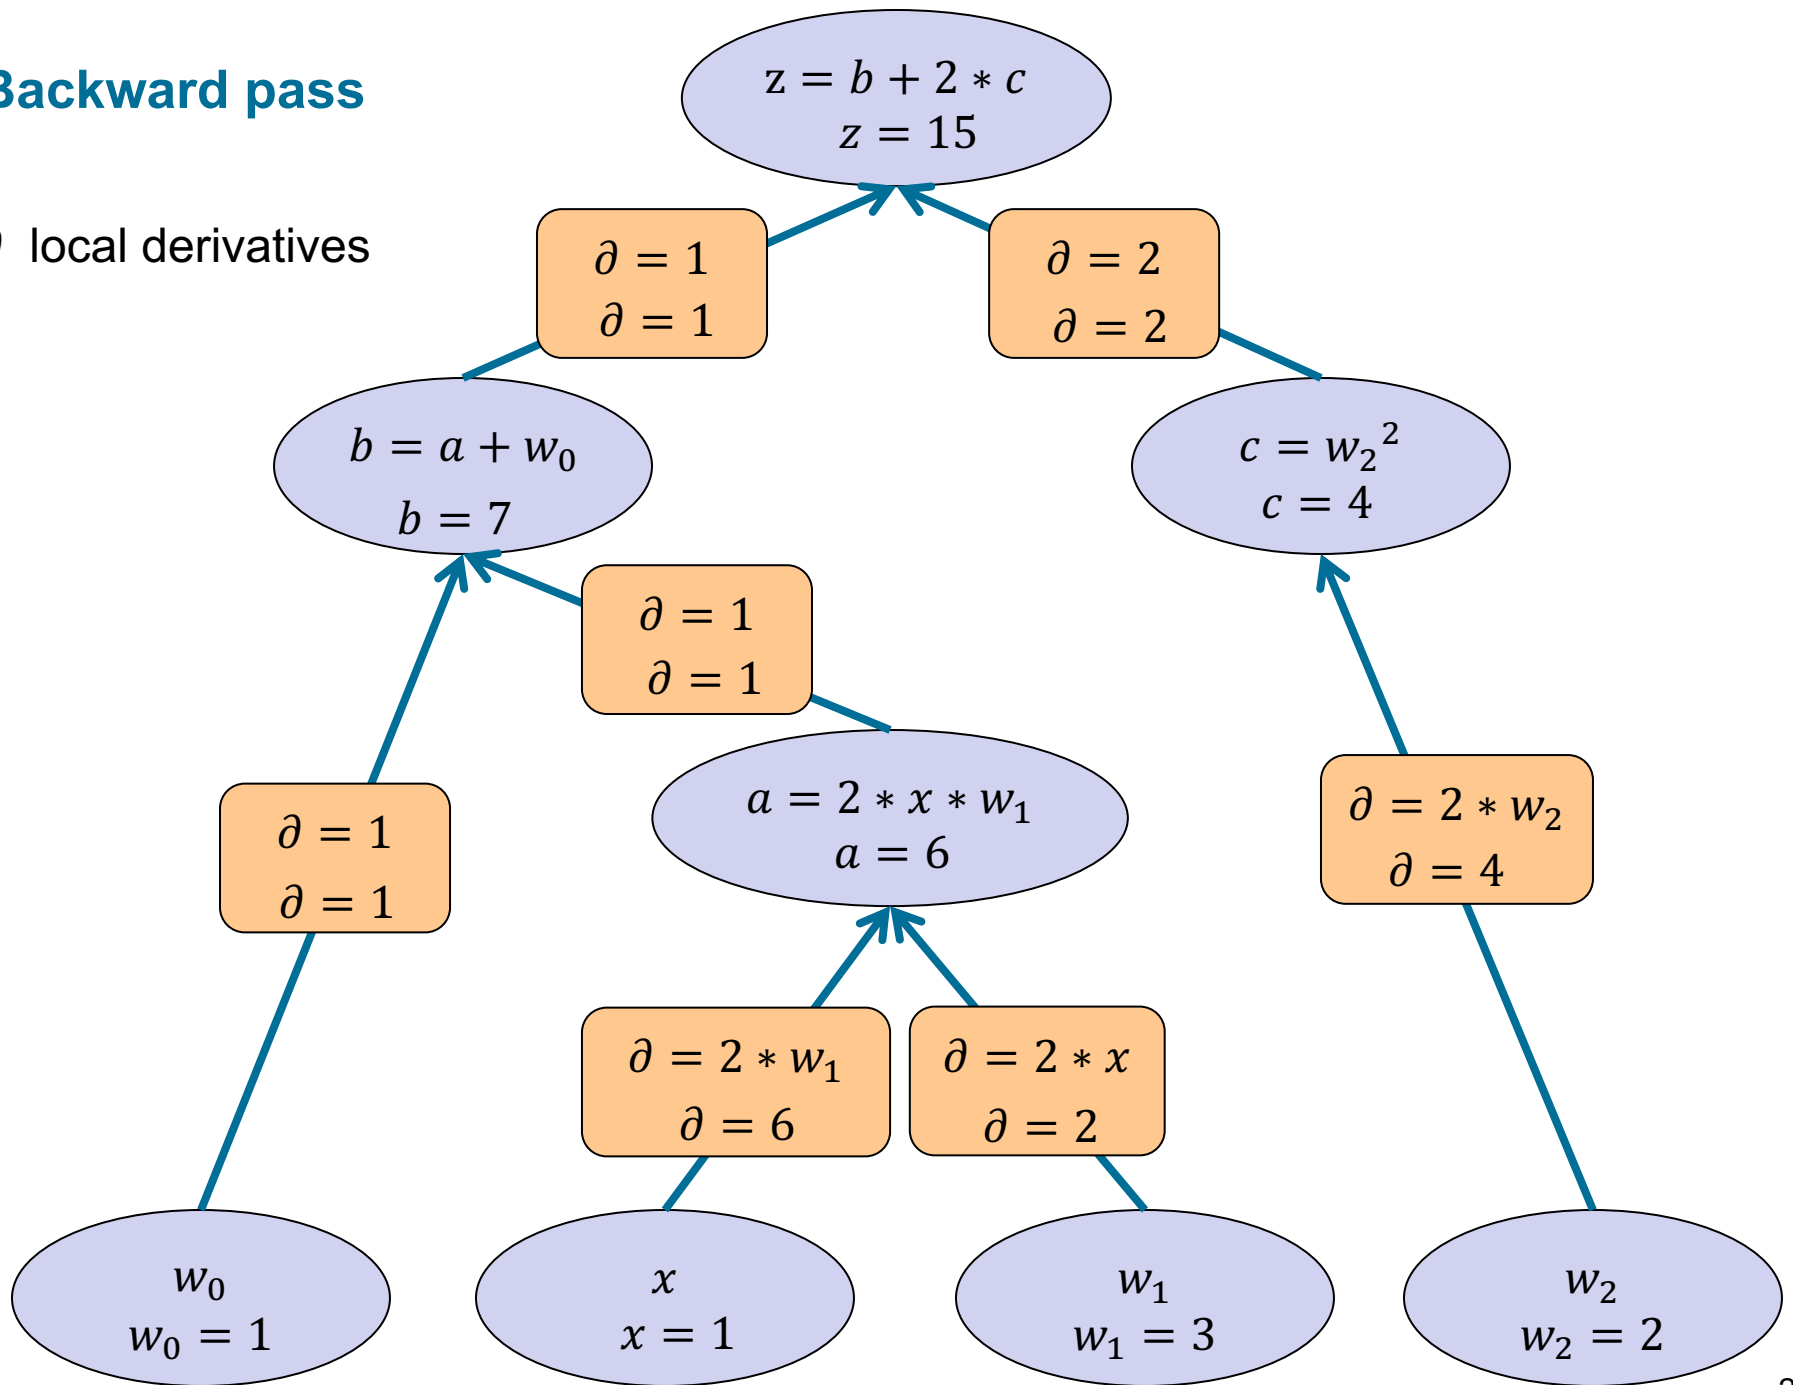
Forward pass

∂ local derivatives



Backward pass

∂ local derivatives



Gradient & Chain rule

- We need the gradient of z regarding \mathbf{w} for optimization

$$\nabla_{\mathbf{w}} z = \begin{bmatrix} \frac{\partial z}{\partial w_0} & \frac{\partial z}{\partial w_1} & \frac{\partial z}{\partial w_2} \end{bmatrix}$$

- We calculate it using chain rule and local derivatives:

$$\frac{\partial z}{\partial w_0} = \frac{\partial z}{\partial b} \frac{\partial b}{\partial w_0}$$

$$\frac{\partial z}{\partial w_1} = \frac{\partial z}{\partial b} \frac{\partial b}{\partial a} \frac{\partial a}{\partial w_1}$$

$$\frac{\partial z}{\partial w_2} = \frac{\partial z}{\partial c} \frac{\partial c}{\partial w_2}$$

Backpropagation

$$\frac{\partial z}{\partial w_0} = \frac{\partial z}{\partial b} \frac{\partial b}{\partial w_0} = 1 * 1 = 1$$

$$\frac{\partial z}{\partial w_1} = \frac{\partial z}{\partial b} \frac{\partial b}{\partial a} \frac{\partial a}{\partial w_1} = 1 * 1 * 2 = 2$$

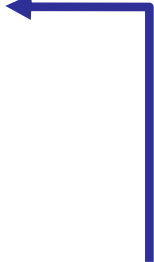
$$\frac{\partial z}{\partial w_2} = \frac{\partial z}{\partial c} \frac{\partial c}{\partial w_2} = 2 * 4 = 8$$

Agenda

- Introduction
- Non-linearities
- Forward pass & backpropagation
- **Softmax & loss function**
- Optimization & regularization

Softmax

- Given the output vector \mathbf{z} of a neural networks model with K output classes
- softmax turns the vector to a probability distribution

$$\text{softmax}(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$



normalization term

Softmax – numeric example

- $K = 4$ classes

$$\mathbf{z} = \begin{bmatrix} 1 \\ 2 \\ 5 \\ 6 \end{bmatrix}$$

$$\text{softmax}(\mathbf{z}) = \begin{bmatrix} 0.004 \\ 0.013 \\ 0.264 \\ 0.717 \end{bmatrix}$$



Softmax characteristics

- The exponential function in softmax makes the highest value becomes separated from the others
- Softmax identifies the “*max*” but in a “*soft*” way!
- Softmax makes competition between the predicted output values, so that in the extreme case, “*winner takes all*”
 - Winner-takes-all: one output is 1 and the rest are 0
 - This resembles the competition between nearby neurons in the cortex

Negative Log Likelihood (NLL) Loss

- NLL loss function is commonly used in neural networks to optimize classification tasks:

$$\mathcal{L} = -\mathbb{E}_{\mathbf{x}, y \sim \mathcal{D}} \log p(y|\mathbf{x}; \mathbb{W})$$

- \mathcal{D} the set of (training) data
 - \mathbf{x} input vector
 - y correct output class
- NLL is a form of cross entropy loss

NLL + Softmax

- The choice of output function (such as softmax) is highly related to the selection of loss function. These two should fit with each other!
- Softmax and NLL are a *good* pair
- To see why, let's calculate the final NLL loss function when softmax is used at output layer (next page)

NLL + Softmax

- Loss function for one data point: $\mathcal{L}(f(\mathbf{x}; \mathbf{w}), y)$
- \mathbf{z} the output vector of network before applying softmax
- y the index of the correct class

$$\begin{aligned}\mathcal{L}(f(\mathbf{x}; \mathbf{w}), y) &= -\log p(y|\mathbf{x}; \mathbb{W}) \\ &= -\log \frac{e^{z_y}}{\sum_{j=1}^K e^{z_j}} \\ &= -z_y + \log \sum_{j=1}^K e^{z_j}\end{aligned}$$

NLL + Softmax – example 2

$$\mathbf{z} = \begin{bmatrix} 1 \\ 2 \\ 0.5 \\ 6 \end{bmatrix}$$

- If the correct class is the first one, $y = 0$:

$$\mathcal{L} = -1 + \log(e^1 + e^2 + e^{0.5} + e^6) = -1 + 6.02 = \mathbf{5.02}$$

- If the correct class is the third one, $y = 2$:

$$\mathcal{L} = -0.5 + \log(e^1 + e^2 + e^{0.5} + e^6) = -0.5 + 6.02 = \mathbf{5.52}$$

- If the correct class is the fourth one, $y = 3$:

$$\mathcal{L} = -6 + \log(e^1 + e^2 + e^{0.5} + e^6) = -6 + 6.02 = \mathbf{0.02}$$

NLL + Softmax – example 1

$$\mathbf{z} = \begin{bmatrix} 1 \\ 2 \\ 5 \\ 6 \end{bmatrix}$$

- If the correct class is the first one, $y = 0$:

$$\mathcal{L} = -1 + \log(e^1 + e^2 + e^5 + e^6) = -1 + 6.33 = \mathbf{5.33}$$

- If the correct class is the third one, $y = 2$:

$$\mathcal{L} = -5 + \log(e^1 + e^2 + e^5 + e^6) = -5 + 6.33 = \mathbf{1.33}$$

- If the correct class is the fourth one, $y = 3$:


$$\mathcal{L} = -6 + \log(e^1 + e^2 + e^5 + e^6) = -6 + 6.33 = \mathbf{0.33}$$

Agenda

- Introduction
- Non-linearities
- Forward pass & backpropagation
- Softmax & loss function
- **Optimization & regularization**

Stochastic Gradient Descent (SGD)

- For every $w \in \mathbb{W}$ and for m training data points



Stochastic Gradient Descent algorithm

- A set of parameters \mathbf{w}
- A **learning rate** η
- Loop until some exit criteria are met
 - Sample a **minibatch** of m data points from \mathcal{D}
 - Compute gradient (vectors) of parameters:

$$\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\mathbf{w}} \sum_i \mathcal{L}(f(\mathbf{x}^{(i)}; \mathbf{w}), y^{(i)})$$

- Update the parameters by taking a step in the opposite direction of the corresponding gradients:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \mathbf{g}$$

- Reduce learning rate (**annealing**) if some criteria are met or based on a schedule

Sampling size

- If only one data point is used in every step; $m = 1$
 - Fast
 - learns **online**
 - Training can become unstable with a lot of fluctuations
- If all data points are used in every step; $m = N$
 - Also called **Batch Gradient Descent**
 - Training can take very long time
- If m is between these
 - Also called **Mini-Batch Gradient Descent**
 - Typical setting for training deep learning models

Other gradient-based optimizations


- Limitations of the mentioned SGD algorithms
 - Choosing learning rate is hard
 - Choosing annealing method/rate is hard
 - Same learning rate is applied to all parameters
 - Can get trapped in non-optimal local minima and saddle points
- Some other commonly used algorithms:
 - Nesterov accelerated gradient
 - Adagrad
 - Adam

Regularization techniques for neural networks and deep learning

- Parameter norm penalties (discussed in previous lecture)
- **Early stopping**
- **Dropout**
- Batch normalization
- Transfer learning
- Multitask learning
- Unsupervised / Semi-supervised pre-training
- Noise robustness
- Dataset augmentation
- Ensemble
- Adversarial training


Early Stopping

- Run the model for several steps (epochs), and in each step evaluate the model on the validation set
- Store the model if the evaluation results improve
- At the end, take the stored model (with best validation results) as the final model




Dropout

- **Key idea:** prune neural network by removing some hidden units stochastically
- At training time for each data point:
 - Each hidden unit's output is multiplied to zero based on a **dropout probability** (like 0.6)




(a) Standard Neural Net




(b) After applying dropout.

Dropout

- At test time:
 - All hidden units are used
 - The output of each hidden is multiplied to the dropout probability



(a) Standard Neural Net



(b) After applying dropout.

Dropout – characteristics

- Computationally inexpensive but a powerful method
- Dropout can be viewed as a geometric average of an exponential number of networks → Ensemble
- Dropout prevents hidden units from forming co-dependencies amongst each other
- Every hidden unit learns to perform well regardless of other units